

# Carseats R Notebook

[CODE ▾](#)

Cooper M. Stansbury (63919844)

## 1 Introduction

Here is my analysis of `data/Carseats_org.csv`.

### 1.1 Configuration

[HIDE](#)

```
library(leaps)
library(stringr)
library(caret)
library(ggplot2)
library(DataExplorer)
library(dplyr)
library(ggExtra)
library(RColorBrewer)
library(plotly)
library(corrplot)
library(htmltools)
library(MASS)
```

### 1.2 System Information

Due to the large number of libraries in use I have provided system information.

[HIDE](#)

```
sessionInfo()
```

```
R version 3.5.1 (2018-07-02)
Platform: x86_64-apple-darwin13.4.0 (64-bit)
Running under: macOS High Sierra 10.13.6

Matrix products: default
BLAS: /System/Library/Frameworks/Accelerate.framework/Versions/A/Frameworks/vecLib.framework/Versions/A/libBLAS.dylib
LAPACK: /anaconda3/lib/R/lib/libRblas.dylib

locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

attached base packages:
[1] stats      graphics  grDevices  utils      datasets  methods   base

other attached packages:
[1] htmltools_0.3.6      corrplot_0.84        plotly_4.8.0         RColorBrewer_1.1-2   ggExtra_0.8
[6] dplyr_0.8.0.1        DataExplorer_0.7.0   mgcv_1.8-26          nlme_3.1-137         caret_6.0-81
[11] ggplot2_3.1.0        lattice_0.20-38      leaps_3.0

loaded via a namespace (and not attached):
[1] Rcpp_1.0.0           lubridate_1.7.4      tidyr_0.8.2          class_7.3-15         assertthat_0.2.0
[6] digest_0.6.18        ipred_0.9-8          foreach_1.4.4        mime_0.6              R6_2.4.0
[11] plyr_1.8.4           stats4_3.5.1         evaluate_0.12         httr_1.4.0           pillar_1.3.1
[16] rlang_0.3.1          lazyeval_0.2.1       rstudioapi_0.9.0     data.table_1.12.0    miniUI_0.1.1.1
[21] rpart_4.1-13         Matrix_1.2-15        rmarkdown_1.11       splines_3.5.1         gower_0.1.2
[26] stringr_1.4.0        htmlwidgets_1.3      igraph_1.2.4         munsell_0.5.0        shiny_1.2.0
[31] httpuv_1.4.5.1       compiler_3.5.1       xfun_0.4             pkgconfig_2.0.2      nnet_7.3-12
[36] tidyselect_0.2.5     tibble_2.0.1         gridExtra_2.3        prodlim_2018.04.18   codetools_0.2-16
[41] viridisLite_0.3.0    later_0.8.0          crayon_1.3.4         withr_2.1.2          MASS_7.3-51.1
[46] recipes_0.1.4        ModelMetrics_1.1.0   grid_3.5.1           xtable_1.8-3         jsonlite_1.6
[51] gtable_0.2.0         magrittr_1.5         scales_1.0.0         stringi_1.3.1        reshape2_1.4.3
[56] promises_1.0.1       timeDate_3043.102    generics_0.0.2       lava_1.6.5           iterators_1.0.10
[61] tools_3.5.1          glue_1.3.0           purrr_0.2.5          networkD3_0.4        parallel_3.5.1
[66] survival_2.43-3      yaml_2.2.0           colorspace_1.4-0     knitr_1.21
```

HIDE

```
sapply(c('repr', 'IRdisplay', 'IRkernel'), function(p) paste(packageVersion(p)))
```

```
repr IRdisplay IRkernel
"0.19.2" "0.7.0" "0.8.15"
```

## 2 Data

First we inspect the raw records using `bash`. Always a good idea to look at things before you mash 'em into an IDE.

```
head -n 5 data/Carseats_org.csv
```

```
"", "Sales", "CompPrice", "Income", "Advertising", "Population", "Price", "ShelveLoc", "Age", "Education", "Urban", "US"
"1", 9.5, 138, 73, 11, 276, 120, "Bad", 42, 17, "Yes", "Yes"
"2", 11.22, 111, 48, 16, 260, 83, "Good", 65, 10, "Yes", "Yes"
"3", 10.06, 113, 35, 10, 269, 80, "Medium", 59, 12, "Yes", "Yes"
"4", 7.4, 117, 100, 4, 466, 97, "Medium", 55, 14, "Yes", "Yes"
```

Now I load the data into `r`, and drop the "ID" column.

```
carseats <- read.csv("data/Carseats_org.csv", header = T, stringsAsFactors = T)
drops <- c("X")
carseats <- carseats[, !(names(carseats) %in% drops)]
head(carseats, 10)
```

Here I create two new data frames to manage numeric and categorical data.

HIDE

HIDE

```
# get vectors of continuous and categorical cols
nums <- dplyr::select_if(carseats, is.numeric)
cats <- dplyr::select_if(carseats, is.factor)
nums[sample(nrow(nums), 10), ]
```

HIDE

```
cats[sample(nrow(cats), 10), ]
```

Let's get some quick summaries of each:

HIDE

```
print('Numeric Summaries')
```

```
[1] "Numeric Summaries"
```

HIDE

```
summary(nums)
```

Sales	CompPrice	Income	Advertising	Population	Price
Min. : 0.000	Min. : 77	Min. : 21.00	Min. : 0.000	Min. : 10.0	Min. : 24.0
1st Qu.: 5.390	1st Qu.:115	1st Qu.: 42.75	1st Qu.: 0.000	1st Qu.:139.0	1st Qu.:100.0
Median : 7.490	Median :125	Median : 69.00	Median : 5.000	Median :272.0	Median :117.0
Mean : 7.496	Mean :125	Mean : 68.66	Mean : 6.635	Mean :264.8	Mean :115.8
3rd Qu.: 9.320	3rd Qu.:135	3rd Qu.: 91.00	3rd Qu.:12.000	3rd Qu.:398.5	3rd Qu.:131.0
Max. :16.270	Max. :175	Max. :120.00	Max. :29.000	Max. :509.0	Max. :191.0

Age	Education
Min. :25.00	Min. :10.0
1st Qu.:39.75	1st Qu.:12.0
Median :54.50	Median :14.0
Mean :53.32	Mean :13.9
3rd Qu.:66.00	3rd Qu.:16.0
Max. :80.00	Max. :18.0

HIDE

```
print('Categorical Summaries')
```

```
[1] "Categorical Summaries"
```

HIDE

```
summary(cats)
```

ShelveLoc	Urban	US
Bad : 96	No :118	No :142
Good : 85	Yes:282	Yes:258
Medium:219		

HIDE

```
str(nums)
```

```
'data.frame':   400 obs. of  8 variables:
 $ Sales      : num  9.5 11.22 10.06 7.4 4.15 ...
 $ CompPrice  : int  138 111 113 117 141 124 115 136 132 132 ...
 $ Income     : int  73 48 35 100 64 113 105 81 110 113 ...
 $ Advertising: int  11 16 10 4 3 13 0 15 0 0 ...
 $ Population : int  276 260 269 466 340 501 45 425 108 131 ...
 $ Price      : int  120 83 80 97 128 72 108 120 124 124 ...
 $ Age        : int  42 65 59 55 38 78 71 67 76 76 ...
 $ Education  : int  17 10 12 14 13 16 15 10 10 17 ...
```

HIDE

```
str(cats)
```

```
'data.frame':   400 obs. of  3 variables:
 $ ShelfLoc: Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2 3 3 ...
 $ Urban   : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 2 1 1 ...
 $ US      : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...
```

## 2.1 Data Dimensionality

This command is to inspect the different data types in the data.

HIDE

```
str(carseats)
```

```
'data.frame':   400 obs. of  11 variables:
 $ Sales      : num  9.5 11.22 10.06 7.4 4.15 ...
 $ CompPrice  : int  138 111 113 117 141 124 115 136 132 132 ...
 $ Income     : int  73 48 35 100 64 113 105 81 110 113 ...
 $ Advertising: int  11 16 10 4 3 13 0 15 0 0 ...
 $ Population : int  276 260 269 466 340 501 45 425 108 131 ...
 $ Price      : int  120 83 80 97 128 72 108 120 124 124 ...
 $ ShelfLoc   : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2 3 3 ...
 $ Age        : int  42 65 59 55 38 78 71 67 76 76 ...
 $ Education  : int  17 10 12 14 13 16 15 10 10 17 ...
 $ Urban      : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1 ...
 $ US         : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...
```

HIDE

```
print("")
```

```
[1] ""
```

HIDE

```
print(paste('Number of Columns:', ncol(carseats)))
```

```
[1] "Number of Columns: 11"
```

HIDE

```
print(paste('Number of Numeric Columns:', ncol(nums)))
```

```
[1] "Number of Numeric Columns: 8"
```

HIDE

```
print(paste('Number of Categorical Columns:', ncol(cats)))
```

```
[1] "Number of Categorical Columns: 3"
```

HIDE

```
dim(carseats)
```

```
[1] 400 11
```

HIDE

```
dim(nums)
```

```
[1] 400 8
```

HIDE

```
dim(cats)
```

```
[1] 400 3
```

Here's a quick way to examine general properties of the data:

HIDE

```
DataExplorer::introduce(data=carseats)
```

Finally, I want to look at the first and last rows of the data set. Just to be safe:

HIDE

```
head(carseats, 2)
```

HIDE

```
tail(carseats, 2)
```

HIDE

```
head(nums, 2)
```

HIDE

```
tail(nums, 2)
```

HIDE

```
head(cats, 2)
```

HIDE

```
tail(cats, 2)
```

## 3 Numeric Plotting

I start out with a few general scatter plots.

[HIDE](#)

```
plot_ly(data=carseats,  
        x=~Age,  
        y=~Sales,  
        mode = 'markers',  
        type = 'scatter',  
        color=~ShelveLoc) %>%  
  layout(title = "Age, Shelf Location, and Sales Scatter Plot", width=900)
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly()

This plot below shows good separation and a weak linear trend. These variables are worth investigating.

[HIDE](#)

```
plot_ly(data=carseats,  
        x=~Price,  
        y=~Sales,  
        mode = 'markers',  
        type = 'scatter',  
        color=~ShelveLoc) %>%  
  layout(title = "Price, Shelf Location, and Sales Scatter Plot", width=900)
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly()

Here we inspect the density of the 'Price vs Sales' relationship:

HIDE

```
plot_ly(data=carseats,  
        x=~Price,  
        y=~Sales,  
        mode = 'markers',  
        size = ~Price,  
        type = 'scatter',  
        colors = "Dark2",  
        alpha = .6) %>%  
  layout(title = "Price, US, and Sales Scatter Plot", width=900)
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly() `line.width` does not currently support multiple values. `line.width` does not currently support multiple values.

## 4 Normalization

I choose to normalize the numeric data in order to be able to plot each variable on the same scale. This will allow me to investigate the variation of each predictor relative to Sales.

HIDE

```
preObj <- preProcess(nums, method=c("center", "scale"))
scaled nums <- predict(preObj, nums)
head(scaled nums, 2)
```

HIDE

```
tail(scaled nums, 2)
```

HIDE

```
str(scaled nums)
```

```
'data.frame': 400 obs. of 8 variables:
 $ Sales      : num  0.7095 1.3185 0.9078 -0.0341 -1.1849 ...
 $ CompPrice  : num  0.849 -0.911 -0.781 -0.52 1.045 ...
 $ Income     : num  0.155 -0.738 -1.203 1.12 -0.166 ...
 $ Advertising: num  0.656 1.408 0.506 -0.396 -0.547 ...
 $ Population : num  0.0757 -0.0328 0.0282 1.3649 0.51 ...
 $ Price      : num  0.178 -1.385 -1.512 -0.794 0.515 ...
 $ Age        : num  -0.699 0.721 0.35 0.104 -0.946 ...
 $ Education  : num  1.183 -1.4882 -0.725 0.0382 -0.3434 ...
```

HIDE

```
print("")
```

```
[1] ""
```

HIDE

```
summary(scaled nums)
```

Sales	CompPrice	Income	Advertising	Population
Min. : -2.65440	Min. : -3.12856	Min. : -1.70290	Min. : -0.9977	Min. : -1.72918
1st Qu.: -0.74584	1st Qu.: -0.65049	1st Qu.: -0.92573	1st Qu.: -0.9977	1st Qu.: -0.85387
Median : -0.00224	Median : 0.00163	Median : 0.01224	Median : -0.2459	Median : 0.04858
Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000
3rd Qu.: 0.64575	3rd Qu.: 0.65375	3rd Qu.: 0.79834	3rd Qu.: 0.8067	3rd Qu.: 0.90693
Max. : 3.10670	Max. : 3.26225	Max. : 1.83458	Max. : 3.3630	Max. : 1.65671

Price	Age	Education
Min. : -3.87702	Min. : -1.74827	Min. : -1.48825
1st Qu.: -0.66711	1st Qu.: -0.83779	1st Qu.: -0.72504
Median : 0.05089	Median : 0.07268	Median : 0.03816
Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
3rd Qu.: 0.64219	3rd Qu.: 0.78255	3rd Qu.: 0.80137
Max. : 3.17633	Max. : 1.64673	Max. : 1.56457

## 4.1 Distributions

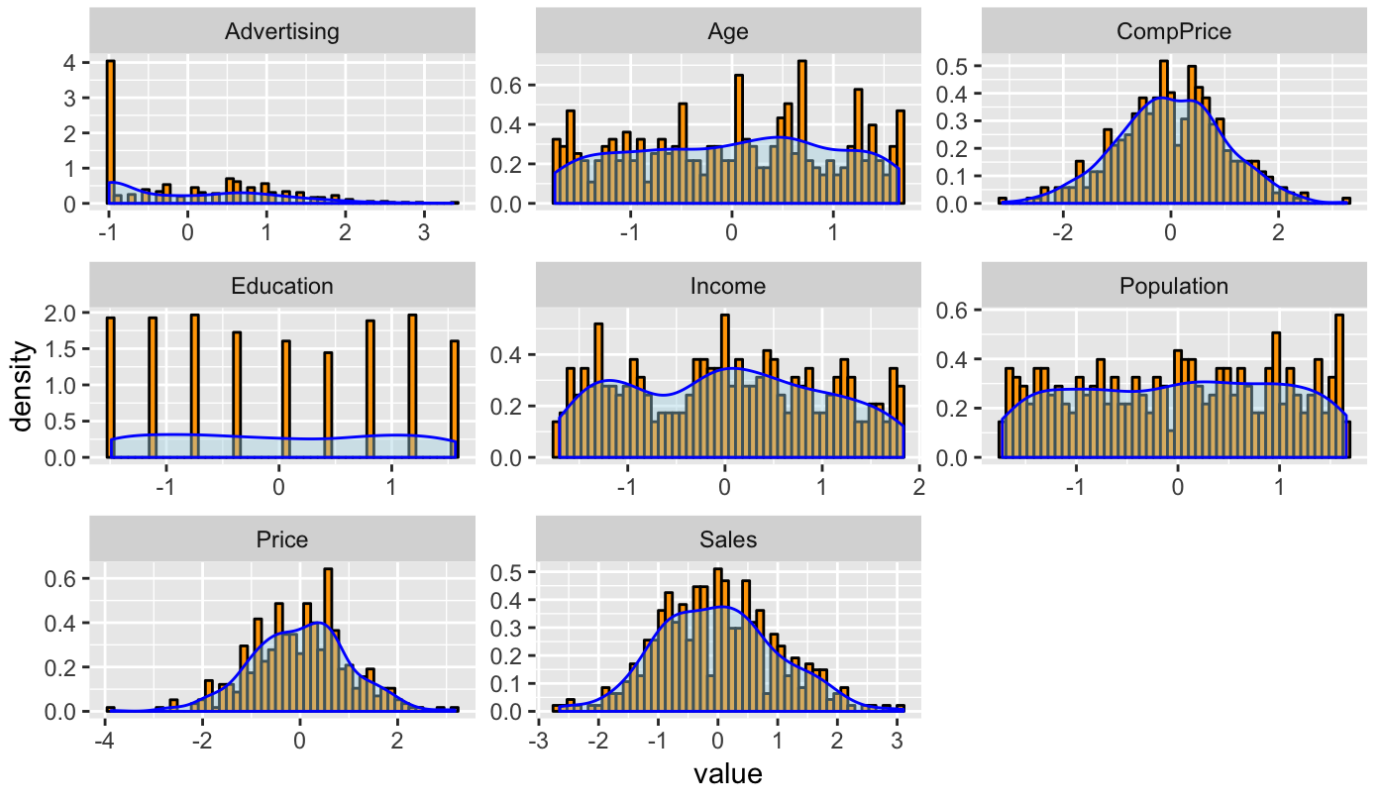
Here are scaled distributions (histograms and density plots) for each numeric variable, including Sales. The variables relating to money (\$) tend to be approximately normal. Many other variables tend to be approximately uniform, which does not bode well for their predictive power.

HIDE

```
scaled nums %>%
  tidyr::gather() %>%
  ggplot(aes(x=value, y=..density..)) +
  ggtitle('Distributions of Continuous Variables (scaled)') +
  facet_wrap(~ key, scales = "free") +
  geom_histogram(fill=I("orange"), col=I("black"), bins = 50) +
  facet_wrap(~ key, scales = "free") +
  geom_density(color="blue", fill='light blue', alpha = 0.4)
```



## Distributions of Continuous Variables (scaled)



Here we plot all numeric variables against their distributions. This is just another way to examine the information shown above.

HIDE

```
scaled.nums %>%
  tidyr::gather() %>%
  plot_ly(x=~key, y=~value,
    type = "box",
    boxpoints = "all",
    jitter = 0.4,
    pointpos = 0,
    color = ~key,
    colors = "Dark2") %>%
  subplot(shareX = TRUE) %>%
  layout(title = "Numeric Variable Distributions (scaled)",
    yaxis=list(title='Standard Deviation'),
    xaxis=list(title='Variable'),
    autosize=FALSE,
    width=900,
    height=500)
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly()

## 4.2 Scatterplots

Here we plot all numeric variables against Sales (scaled). This allows us to investigate possible linear relationships between that variable and Sales. As shown below, only 'Price' appears to have a linear relationship worth investigating. This took me so long to figure out.

HIDE

```
numeric.scatterplots <- htmtools::tagList()
count = 1
for (i in names(scaled.nums[, -1])) {

  numeric.scatterplots[[count]] <- plot_ly(scaled.nums, x=scaled.nums[, i], y=scaled.nums$Sales,
    colors = 'RdYlGn',
    mode = 'markers',
    type = 'scatter',
    size = scaled.nums$Sales^2,
    color = scaled.nums$Sales,
    marker = list(line = list(color = 'black', width = 2)),
    name=paste(i)) %>%
    layout(title = paste(i, "vs Sales (scaled)<br>Size=Sales^2"),
      yaxis=list(title='Sales'),
      xaxis=list(title=i),
      showlegend = FALSE)

  count = count + 1
}
numeric.scatterplots
```

```
'line.width' does not currently support multiple values.'line.width' does not currently support multiple val  
ues.'line.width' does not currently support multiple values.'line.width' does not currently support multiple  
values.'line.width' does not currently support multiple values.'line.width' does not currently support multi  
ple values.'line.width' does not currently support multiple values.'line.width' does not currently support m  
ultiple values.'line.width' does not currently support multiple values.'line.width' does not currently suppo  
rt multiple values.'line.width' does not currently support multiple values.'line.width' does not currently s  
upport multiple values.'line.width' does not currently support multiple values.'line.width' does not current  
ly support multiple values.
```



By adding naive regression lines to a few scatter plots we can confirm our suspicions:

```

fit.Pop <- lm(Sales ~ Population, data = scaled.nums)
fit.Age <- lm(Sales ~ Age, data = scaled.nums)
fit.CompPrice <- lm(Sales ~ CompPrice, data = scaled.nums)
fit.Price <- lm(Sales ~ Price, data = scaled.nums)
regression.scatterplots <- htmltools::tagList()
regression.scatterplots[[1]] <- plot_ly(scaled.nums,
  x = ~Population,
  name = 'Population vs Sales Regression Line') %>%
  add_markers(y = ~Sales,
    name = 'Population vs Sales Observations') %>%
    add_lines(x = ~Population,
      y = fitted(fit.Pop)) %>%
      layout(title = "Population vs Sales",
        yaxis=list(title='Sales',
          xaxis=list(title='Population')),
        showlegend = FALSE)

regression.scatterplots[[2]] <- plot_ly(scaled.nums,
  x = ~Age,
  name = 'Age vs Sales Regression Line') %>%
  add_markers(y = ~Sales,
    name = 'Age vs Sales Observations') %>%
    add_lines(x = ~Age,
      y = fitted(fit.Age)) %>%
      layout(title = "Age vs Sales",
        yaxis=list(title='Sales',
          xaxis=list(title='Age')),
        showlegend = FALSE)

regression.scatterplots[[3]] <- plot_ly(scaled.nums,
  x = ~CompPrice,
  name = 'CompPrice vs Sales Regression Line') %>%
  add_markers(y = ~Sales,
    name = 'CompPrice vs Sales Observations') %>%
    add_lines(x = ~CompPrice,
      y = fitted(fit.CompPrice)) %>%
      layout(title = "CompPrice vs Sales",
        yaxis=list(title='Sales',
          xaxis=list(title='CompPrice')),
        showlegend = FALSE)

regression.scatterplots[[4]] <- plot_ly(scaled.nums,
  x = ~Price,
  name = 'Price vs Sales Regression Line') %>%
  add_markers(y = ~Sales,
    name = 'Price vs Sales Observations') %>%
    add_lines(x = ~Price,
      y = fitted(fit.Price)) %>%
      layout(title = "Price vs Sales",
        yaxis=list(title='Sales',
          xaxis=list(title='Price')),
        showlegend = FALSE)

regression.scatterplots

```



Let's compare the slopes:

HIDE

```
scaled.nums %>%
  plot_ly(y = ~Sales) %>%
    add_lines(x= ~Population, y = fitted(fit.Pop),
              name = "fit.Pop slope", line = list(shape = "linear")) %>%
    add_lines(x= ~Age, y = fitted(fit.Age),
              name = "fit.Age slope", line = list(shape = "linear")) %>%
    add_lines(x= ~CompPrice, y = fitted(fit.CompPrice),
              name = "fit.CompPrice slope", line = list(shape = "linear")) %>%
    add_lines(x= ~Price, y = fitted(fit.Price),
              name = "fit.Price slope", line = list(shape = "linear")) %>%

  layout(title = "Regression Lines vs Sales (scaled)",
         autosize=FALSE,
         width=900,
         yaxis=list(title='Sales'),
         xaxis=list(title='Scaled Numeric Variable'))
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly()

Here's a pretty graphic that doesn't help me understand anything about the data.

HIDE

```
y = scaled.nums$Sales
x = scaled.nums$Price
s <- subplot(
  plot_ly(x = x, color = I("black"), type = 'histogram'),
  plotly_empty(),
  plot_ly(x = x, y = y, type = 'histogram2dcontour', showscale = F),
  plot_ly(y = y, color = I("black"), type = 'histogram'),
  nrows = 2, heights = c(0.2, 0.8), widths = c(0.8, 0.2),
  shareX = TRUE, shareY = TRUE, titleX = FALSE, titleY = FALSE)
```

No trace type specified and no positional attributes specifiedNo trace type specified:  
Based on info supplied, a 'scatter' trace seems appropriate.  
Read more about this trace type -> <https://plot.ly/r/reference/#scatter>  
No scatter mode specified:  
Setting the mode to markers  
Read more about this attribute -> <https://plot.ly/r/reference/#scatter-mode>

HIDE

```
layout(s, showlegend = FALSE, autosize=FALSE,  
       width=700,  
       height=500,  
       yaxis=list(title='Sales'),  
       xaxis=list(title='Price'))
```

Specifying width/height in layout() is now deprecated.  
Please specify in ggplotly() or plot\_ly()

## 5 Categorical Plotting

First, let's create a data frame that we can use:

HIDE

```
categorical.by.sales = cbind(Sales = scaled.nums$Sales, cats)  
str(categorical.by.sales)
```

```
'data.frame':   400 obs. of  4 variables:  
 $ Sales      : num  0.7095 1.3185 0.9078 -0.0341 -1.1849 ...  
 $ ShelfLoc   : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2 3 3 ...  
 $ Urban      : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1 ...  
 $ US         : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...
```

Here we can see all categorical by Sales. We suspected that 'ShelveLoc' would be important based on one of the early scatter plots. It seems that this is the case.



```
categorical.boxplots <- htmltools::tagList()
count = 1
for (i in names(categorical.by.sales[, -1])) {

  categorical.boxplots[[count]] <- plot_ly(categorical.by.sales, x=categorical.by.sales[, i], y=categorical.by.

    type = "box",
    boxpoints = "all",
    jitter = .2,
    pointpos = 0,
    color = categorical.by.sales[, i],
    colors = 'Set1',
    name = paste(i)) %>%
    layout(title = paste(i, "vs Sales (scaled)",
      showlegend = TRUE,
      yaxis = list(title = 'Sales Standard Deviation'),
      xaxis = list(title = i))

  count = count + 1
}
categorical.boxplots
```

Here's the same thing, but more musically:

HIDE

```
categorical.violins <- htmltools::tagList()
count = 1
for (i in names(categorical.by.sales[, -1])) {

  categorical.violins[[count]] <- plot_ly(categorical.by.sales, x=categorical.by.sales[, i], y=categorical.by.s

    split = categorical.by.sales[, i],
    type = 'violin',
    colors='Set1',
    name=paste(i),
    box = list(visible = TRUE),
    meanline = list(visible = TRUE)) %>%
      layout(xaxis = list(title = "US"),
             yaxis = list(title = "Sales", zeroline = FALSE))

  count = count + 1
}
categorical.violins
```

## 6 Linear Regression

First, let's merge the data set into a single data frame

HIDE

```
scaled.merged <- cbind(categorical.by.sales[,-1], scaled.nums)
str(scaled.merged)
```

```
'data.frame':  400 obs. of  11 variables:
 $ ShelfLoc   : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2 3 3 ...
 $ Urban      : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1 ...
 $ US         : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...
 $ Sales      : num  0.7095 1.3185 0.9078 -0.0341 -1.1849 ...
 $ CompPrice  : num  0.849 -0.911 -0.781 -0.52 1.045 ...
 $ Income     : num  0.155 -0.738 -1.203 1.12 -0.166 ...
 $ Advertising: num  0.656 1.408 0.506 -0.396 -0.547 ...
 $ Population : num  0.0757 -0.0328 0.0282 1.3649 0.51 ...
 $ Price      : num  0.178 -1.385 -1.512 -0.794 0.515 ...
 $ Age        : num  -0.699 0.721 0.35 0.104 -0.946 ...
 $ Education  : num  1.183 -1.4882 -0.725 0.0382 -0.3434 ...
```

HIDE

```
head(nums, 2)
```

HIDE

```
tail(nums, 2)
```

HIDE

```
head(scaled.merged, 2)
```


HIDE

```
tail(scaled.merged, 2)
```

First, let's look at some things that may give us trouble. Luckily it looks like the only serious correlation is with our dependent variable. We'll want to watch the 'Price' vs 'CompPrice' relationship.

HIDE

```
res <- cor(scaled nums)
corrplot(res, type = "upper", order = "hclust",
  tl.col = "black", tl.srt = 45)
```



It appears that residuals are roughly symmetrical around 0. That's strange. Mostly due to a relatively poor overall fit. Note how close to zero most of the coefficient estimates are.

HIDE

```
simple.lm <- lm(Sales~., data=scaled.merged)
simple.summary <- summary(simple.lm)
print(simple.summary)
```

```
Call:
lm(formula = Sales ~ ., data = scaled.merged)

Residuals:
    Min       1Q   Median       3Q      Max
-1.01598 -0.24463  0.00748  0.23496  1.20797

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -0.73292    0.05999  -12.217  < 2e-16 ***
ShelveLocGood    1.71742    0.05422   31.678  < 2e-16 ***
ShelveLocMedium  0.69286    0.04465   15.516  < 2e-16 ***
UrbanYes        0.04351    0.04000    1.088    0.277
USYes          -0.06519    0.05306   -1.229    0.220
CompPrice       0.50397    0.02252   22.378  < 2e-16 ***
Income         0.15660    0.01828    8.565 2.58e-16 ***
Advertising     0.28987    0.02619   11.066  < 2e-16 ***
Population      0.01085    0.01933    0.561    0.575
Price          -0.79946    0.02239  -35.700  < 2e-16 ***
Age            -0.26413    0.01825  -14.472  < 2e-16 ***
Education      -0.01958    0.01830   -1.070    0.285
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3608 on 388 degrees of freedom
Multiple R-squared:  0.8734,    Adjusted R-squared:  0.8698
F-statistic: 243.4 on 11 and 388 DF,  p-value: < 2.2e-16
```

## 6.1 Linear Models and Subsets

Let's do the same thing, but control the subsets using `leaps`

HIDE

```
regfit.full=regsubsets(Sales~., data=scaled.merged, nvmax=5)
reg.summary=summary(regfit.full)
print(reg.summary)
```

```
Subset selection object
Call: regsubsets.formula(Sales ~ ., data = scaled.merged, nvmax = 5)
11 Variables (and intercept)
      Forced in Forced out
ShelveLocGood    FALSE    FALSE
ShelveLocMedium  FALSE    FALSE
UrbanYes         FALSE    FALSE
USYes           FALSE    FALSE
CompPrice        FALSE    FALSE
Income           FALSE    FALSE
Advertising       FALSE    FALSE
Population        FALSE    FALSE
Price            FALSE    FALSE
Age              FALSE    FALSE
Education         FALSE    FALSE
1 subsets of each size up to 5
Selection Algorithm: exhaustive
      ShelveLocGood ShelveLocMedium UrbanYes USYes CompPrice Income Advertising Population Price Age
1 ( 1 ) " " " " " " " " " " " "
2 ( 1 ) " " " " " " " " " " " "
3 ( 1 ) " " " " " " " " " " " "
4 ( 1 ) " " " " " " " " " " " "
5 ( 1 ) " " " " " " " " " " " "
      Education
1 ( 1 ) " "
2 ( 1 ) " "
3 ( 1 ) " "
4 ( 1 ) " "
5 ( 1 ) " "
```





We'll just take code straight from the example on Canvas...

[HIDE](#)


```
par(mfrow=c(2,2))
plot(reg.summary$rss,xlab="Number of Variables",ylab="RSS",type="l")
plot(reg.summary$adjr2,xlab="Number of Variables",ylab="Adjusted RSq",type="l")
```

[HIDE](#)

```
plot(reg.summary$cp,xlab="Number of Variables",ylab="Cp",type='l')
plot(reg.summary$bic,xlab="Number of Variables",ylab="BIC",type='l')
```


[HIDE](#)

```
plot(regfit.full,scale="r2")
```




HIDE

```
plot(regfit.full, scale="adjr2")
```




HIDE

```
plot(regfit.full, scale="Cp")
```



HIDE

```
plot(regfit.full, scale="bic")
```



## 7 Interaction Terms

Here we define a new model with some interaction terms: a. Income and Advertising b. Income and CompPrice c. Price and Age



```
interaction.lm <- lm(Sales~. + Income*Advertising + Income*CompPrice + Price*Age, data=scaled.merged)
interaction.summary <- summary(interaction.lm)
print(interaction.summary)
```

Call:

```
lm(formula = Sales ~ . + Income * Advertising + Income * CompPrice +
    Price * Age, data = scaled.merged)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-1.04967 -0.23955 -0.00936  0.24591  1.19774
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.730051	0.059236	-12.324	< 2e-16 ***
ShelveLocGood	1.714324	0.053519	32.032	< 2e-16 ***
ShelveLocMedium	0.680563	0.044177	15.405	< 2e-16 ***
UrbanYes	0.045257	0.039380	1.149	0.25117
USYes	-0.069622	0.052329	-1.330	0.18415
CompPrice	0.507567	0.022131	22.935	< 2e-16 ***
Income	0.152164	0.018025	8.442	6.43e-16 ***
Advertising	0.289455	0.025756	11.238	< 2e-16 ***
Population	0.007224	0.018985	0.381	0.70377
Price	-0.797242	0.022003	-36.233	< 2e-16 ***
Age	-0.260783	0.017978	-14.505	< 2e-16 ***
Education	-0.024248	0.018063	-1.342	0.18027
Income:Advertising	0.046378	0.018170	2.552	0.01108 *
CompPrice:Income	-0.059184	0.018911	-3.130	0.00188 **
Price:Age	0.013171	0.017912	0.735	0.46259

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.3538 on 385 degrees of freedom

Multiple R-squared: 0.8792, Adjusted R-squared: 0.8748

F-statistic: 200.1 on 14 and 385 DF, p-value: < 2.2e-16

```
interaction.lm.subsets <- regsubsets(Sales~. + Income*Advertising + Income*CompPrice + Price*Age,
                                     data=scaled.merged, nvmax=5)
interaction.subsets.summary <- summary(interaction.lm.subsets)
print(interaction.subsets.summary)
```

```
Subset selection object
Call: regsubsets.formula(Sales ~ . + Income * Advertising + Income *
  CompPrice + Price * Age, data = scaled.merged, nvmax = 5)
14 Variables (and intercept)
      Forced in Forced out
ShelveLocGood      FALSE      FALSE
ShelveLocMedium    FALSE      FALSE
UrbanYes           FALSE      FALSE
USYes              FALSE      FALSE
CompPrice           FALSE      FALSE
Income             FALSE      FALSE
Advertising         FALSE      FALSE
Population          FALSE      FALSE
Price              FALSE      FALSE
Age                FALSE      FALSE
Education           FALSE      FALSE
Income:Advertising  FALSE      FALSE
CompPrice:Income    FALSE      FALSE
Price:Age           FALSE      FALSE
1 subsets of each size up to 5
Selection Algorithm: exhaustive
```

		ShelveLocGood	ShelveLocMedium	UrbanYes	USYes	CompPrice	Income	Advertising	Population	Price	Age
1	( 1 )	"*"	" "	" "	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	"*"	" "	" "	" "	" "	" "	" "	" "	"*"	" "
3	( 1 )	"*"	" "	" "	" "	"*"	" "	" "	" "	"*"	" "
4	( 1 )	"*"	" "	" "	" "	"*"	" "	"*"	" "	"*"	" "
5	( 1 )	"*"	"*"	" "	" "	"*"	" "	"*"	" "	"*"	" "

		Education	Income:Advertising	CompPrice:Income	Price:Age
1	( 1 )	" "	" "	" "	" "
2	( 1 )	" "	" "	" "	" "
3	( 1 )	" "	" "	" "	" "
4	( 1 )	" "	" "	" "	" "
5	( 1 )	" "	" "	" "	" "

## 7.1 Variable Significance

Below we print the coefficients for the 5th model using the default model selection criteria. All coefficients are relatively small, as we would expect from the EDA above. This pretty much confirms what I would have guessed by looking at the data against sales. We still want to watch out for confounding between 'Price' and 'CompPrice.'

```
coef(interaction.lm.subsets, 5)
```

(Intercept)	ShelveLocGood	ShelveLocMedium	CompPrice	Advertising	Price
-0.6995348	1.6746198	0.6277225	0.5090177	0.2832405	-0.7846476

## 7.2 Second Interaction Model

First, drop columns unneeded from analysis:

```
scaled.merged.slim <- scaled.merged[ , -which(names(scaled.merged) %in% c("US", "Urban"))]
```

A few hyper parameters we'd like to be consistent for all models

```
nvmax <- 3
```

## 7.3 Forward Selection:

```
interaction.subset.fwd <- regsubsets(Sales~. + Income*Advertising + Income*CompPrice + Income*Age,
                                   data=scaled.merged.slim,
                                   nvmax = nvmax,
                                   method="forward")
fwd.subset.summary <- summary(interaction.subset.fwd)
coef(interaction.subset.fwd, 1:nvmax)
```

```
[[1]]
(Intercept) ShelveLocGood
-0.259671      1.221981

[[2]]
(Intercept) ShelveLocGood      Price
-0.2708256      1.2744733      -0.4688868

[[3]]
(Intercept) ShelveLocGood      CompPrice      Price
-0.2709362      1.2749938      0.4932455      -0.7573701
```

## 7.4 Backward Selection:

This is really strange. I can't seem to find any documentation about this, but it appears that this model is actually 'forward.'

```
interaction.subset.bk <- regsubsets(Sales~. + Income*Advertising + Income*CompPrice + Income*Age,
                                   data=scaled.merged.slim,
                                   nvmax = nvmax,
                                   method="backward")
bk.subset.summary <- summary(interaction.subset.bk)
coef(interaction.subset.bk, 1:nvmax)
```

```
[[1]]
(Intercept) ShelveLocGood
-0.259671      1.221981

[[2]]
(Intercept) ShelveLocGood      Price
-0.2708256      1.2744733      -0.4688868

[[3]]
(Intercept) ShelveLocGood      CompPrice      Price
-0.2709362      1.2749938      0.4932455      -0.7573701
```

## 7.5 Exhasutive

```
interaction.subset.ex <- regsubsets(Sales~. + Income*Advertising + Income*CompPrice + Income*Age,
                                   data=scaled.merged.slim,
                                   nvmax = nvmax,
                                   method="exhaustive")
ex.subset.summary <- summary(interaction.subset.ex)
coef(interaction.subset.ex, 1:nvmax)
```

```
[[1]]
(Intercept) ShelfLocGood
-0.259671    1.221981

[[2]]
(Intercept) ShelfLocGood      Price
-0.2708256    1.2744733    -0.4688868

[[3]]
(Intercept) ShelfLocGood    CompPrice      Price
-0.2709362    1.2749938    0.4932455    -0.7573701
```

This is a list that may come in handy.

HIDE

```
model.list = list(list("Forward", interaction.subset.fwd, fwd.subset.summary),
                  list("Backward", interaction.subset.bk, bk.subset.summary),
                  list("Exhaustive", interaction.subset.ex, ex.subset.summary))
```

## 8 Evaluation Metric Plotting





It is interesting to see that each model selected the same variables, in the same order.

HIDE

```
par(mfrow=c(2,2))
plot(fwd.subset.summary$rss,xlab="Number of Variables (forward)",ylab="RSS",type="l")
plot(fwd.subset.summary$adjr2,xlab="Number of Variables (forward)",ylab="Adjusted RSq",type="l")
```

HIDE

```
plot(fwd.subset.summary$cpl,xlab="Number of Variables (forward)",ylab="Cp",type='l')
plot(fwd.subset.summary$bic,xlab="Number of Variables (forward)",ylab="BIC",type='l')
```







HIDE

```
par(mfrow=c(2,2))
plot(bk.subset.summary$rss,xlab="Number of Variables (backward)",ylab="RSS",type="l")
plot(bk.subset.summary$adjr2,xlab="Number of Variables (backward)",ylab="Adjusted RSq",type="l")
```

HIDE

```
plot(bk.subset.summary$cp,xlab="Number of Variables (backward)",ylab="Cp",type='l')
plot(bk.subset.summary$bic,xlab="Number of Variables (backward)",ylab="BIC",type='l')
```







HIDE

```
par(mfrow=c(2,2))
plot(ex.subset.summary$rss,xlab="Number of Variables (exhaustive)",ylab="RSS",type="l")
plot(ex.subset.summary$adjr2,xlab="Number of Variables (exhaustive)",ylab="Adjusted RSq",type="l")
```

HIDE

```
plot(ex.subset.summary$cp,xlab="Number of Variables (exhaustive)",ylab="Cp",type='l')
plot(ex.subset.summary$bic,xlab="Number of Variables (exhaustive)",ylab="BIC",type='l')
```



## 9 Model Equations

Here we print the final equations for each model. Not, they are all the same.

```
for (mod.obj in model.list) {
  mod.name <- mod.obj[[1]]
  best.bic <- min(mod.obj[[3]]$bic)
  mod.num <- which.min(mod.obj[[3]]$bic)
  mod.cc <- coef(mod.obj[[2]], mod.num)

  mod.equation.format <- paste("Y =", paste(round(mod.cc[1], 2),
      paste(round(mod.cc[-1], 2),
        names(mod.cc[-1]),
        sep=" * ", collapse=" + "),
      sep=" + "), "+ e")

  print(paste("Model Selection Method: ", mod.name))
  print(paste("Max BIC:", best.bic))
  print(paste("Model Number: ", mod.num))
  print(paste("Model Equation: ", mod.equation.format))
  print("")
}
```

HIDE

```
[1] "Model Selection Method: Forward"
[1] "Max BIC: -373.710213587368"
[1] "Model Number: 3"
[1] "Model Equation:  $Y = -0.27 + 1.27 * ShelfLocGood + 0.49 * CompPrice + -0.76 * Price + e$ "
[1] ""
[1] "Model Selection Method: Backward"
[1] "Max BIC: -373.710213587368"
[1] "Model Number: 3"
[1] "Model Equation:  $Y = -0.27 + 1.27 * ShelfLocGood + 0.49 * CompPrice + -0.76 * Price + e$ "
[1] ""
[1] "Model Selection Method: Exhaustive"
[1] "Max BIC: -373.710213587368"
[1] "Model Number: 3"
[1] "Model Equation:  $Y = -0.27 + 1.27 * ShelfLocGood + 0.49 * CompPrice + -0.76 * Price + e$ "
[1] ""
```